

Performance benchmark report

Jorge Cubero Toribio

November 2025

1 Introduction

This report investigates performance optimization techniques for matrix multiplication. Three implementations are compared:

1. **Basic Pure Python Implementation:** a triple nested loop using lists of lists.
2. **Dense NumPy Multiplication:** highly optimized BLAS-based dense matrix multiplication using the `@` operator.
3. **Sparse SciPy Multiplication:** compressed sparse row (CSR) representation with 99% sparsity.

The main goal is to evaluate the impact of optimizations and sparsity on:

- Execution time,
- Memory usage,
- Scalability with matrix size,
- Performance under different sparsity levels.

2 Methodology

2.1 Implementation Details

- The **Basic** version uses explicit loops, serving as a baseline for complexity.
- The **Dense** version uses NumPy arrays with contiguous memory layouts.
- The **Sparse** version uses SciPy's CSR format, which stores only non-zero values along with their indices.

2.2 Experimental Setup

Parameter	Description
CPU	Local host (Python execution environment)
Software	Python 3.x, NumPy, SciPy
Matrix Sizes	256, 512, 1024, 2048, 4096
Sparsity Levels	99% (for size test), densities 0.001–0.3 for sparsity sweep
Timing	<code>time.perf_counter()</code>
Memory	<code>nbytes</code> for dense, CSR <code>data+indices+indptr</code> for sparse

Table 1: Experimental setup parameters.

3 Results

3.1 Dense vs. Sparse Multiplication (99% sparsity)

Matrix Size	NumPy Dense (s)	SciPy Sparse (s)
256	0.0293	0.00018
512	0.00217	0.00047
1024	0.0143	0.00171
2048	0.1303	0.0132
4096	0.8694	0.0944

Table 2: Execution time comparison for dense and sparse matrices at 99% sparsity.

Matrix Size	Dense Memory (bytes)	Sparse Memory (bytes)
256	524,288	8,888
512	2,097,152	33,504
1024	8,388,608	129,932
2048	33,554,432	511,512
4096	134,217,728	2,029,652

Table 3: Memory usage comparison between dense and sparse matrices.



Figure 1: Execution Time vs. Matrix Size (log scale). NumPy dense vs. SciPy sparse multiplication.




Figure 2: Memory Usage vs. Matrix Size (log scale). Dense matrices scale quadratically, while sparse matrices grow linearly with non-zero entries.

Observations:

- Sparse matrices are 10–100× times faster at high sparsity (99% zeros).
- Memory use is dramatically reduced — at 4096×4096 , sparse uses about 1.5% of dense memory.
- Dense multiplication scales as $O(N^3)$ in time and $O(N^2)$ in memory.

3.2 Effect of Sparsity on Performance (2048×2048)

Density (% non-zero)	Sparsity (% zeros)	Sparse Time (s)	Memory (MB)
0.1	99.9	0.0034	0.0585
1	99	0.0088	0.512
5	95	0.114	2.52
10	90	0.267	5.04
20	80	0.882	10.07
30	70	1.856	15.11

Table 4: Sparse performance vs. density for 2048×2048 matrices. Dense baseline: 0.1992 s.




Figure 3: Execution Time vs. Density for 2048×2048 matrices. Sparse multiplication surpasses dense until roughly 10% density.

Analysis:

- Sparse multiplication is faster than dense up to about 10% density ($\approx 90\%$ sparsity).
- Memory grows linearly with the number of non-zeros.
- Beyond $\sim 10\%$ density, sparse indexing overhead dominates.

3.3 Scalability and Maximum Efficient Size

Method	Max Size Efficiently Tested	Limiting Factor
Basic (Python)	128×128	CPU time ($O(N^3)$)
NumPy Dense	4096×4096	Memory (134 MB per matrix)
SciPy Sparse (99%)	4096×4096	None observed

Table 5: Maximum efficiently handled matrix sizes.

4 Discussion

- Dense multiplication benefits from cache locality and vectorized BLAS operations.
- Sparse multiplication is advantageous when most elements are zero, since the time complexity approximates $O(kN)$, where k is the average non-zero count per row.
- The break-even point between dense and sparse occurs at roughly 10% density.
- For dense workloads, performance becomes memory-bandwidth limited.

5 Conclusions

- Sparse matrices achieve major speed and memory improvements for sparsity levels above 90%.

- Dense NumPy multiplication remains optimal for dense or moderately sized matrices.
- The pure Python baseline is educational but impractical for performance computing.
- Overall, sparse matrix methods are critical for large-scale or memory-constrained applications.

6 Future Work

Future extensions may include:

- Implementing block-based (tiled) and Strassen algorithms.
- Parallelization via Numba or OpenMP.
- GPU acceleration (e.g., CuPy, PyTorch).
- Testing larger matrices ($N > 10,000$) or distributed systems.

All deliverables are located at: https://github.com/JorsuCT/Big_Data/tree/main/IndividualAssignment